

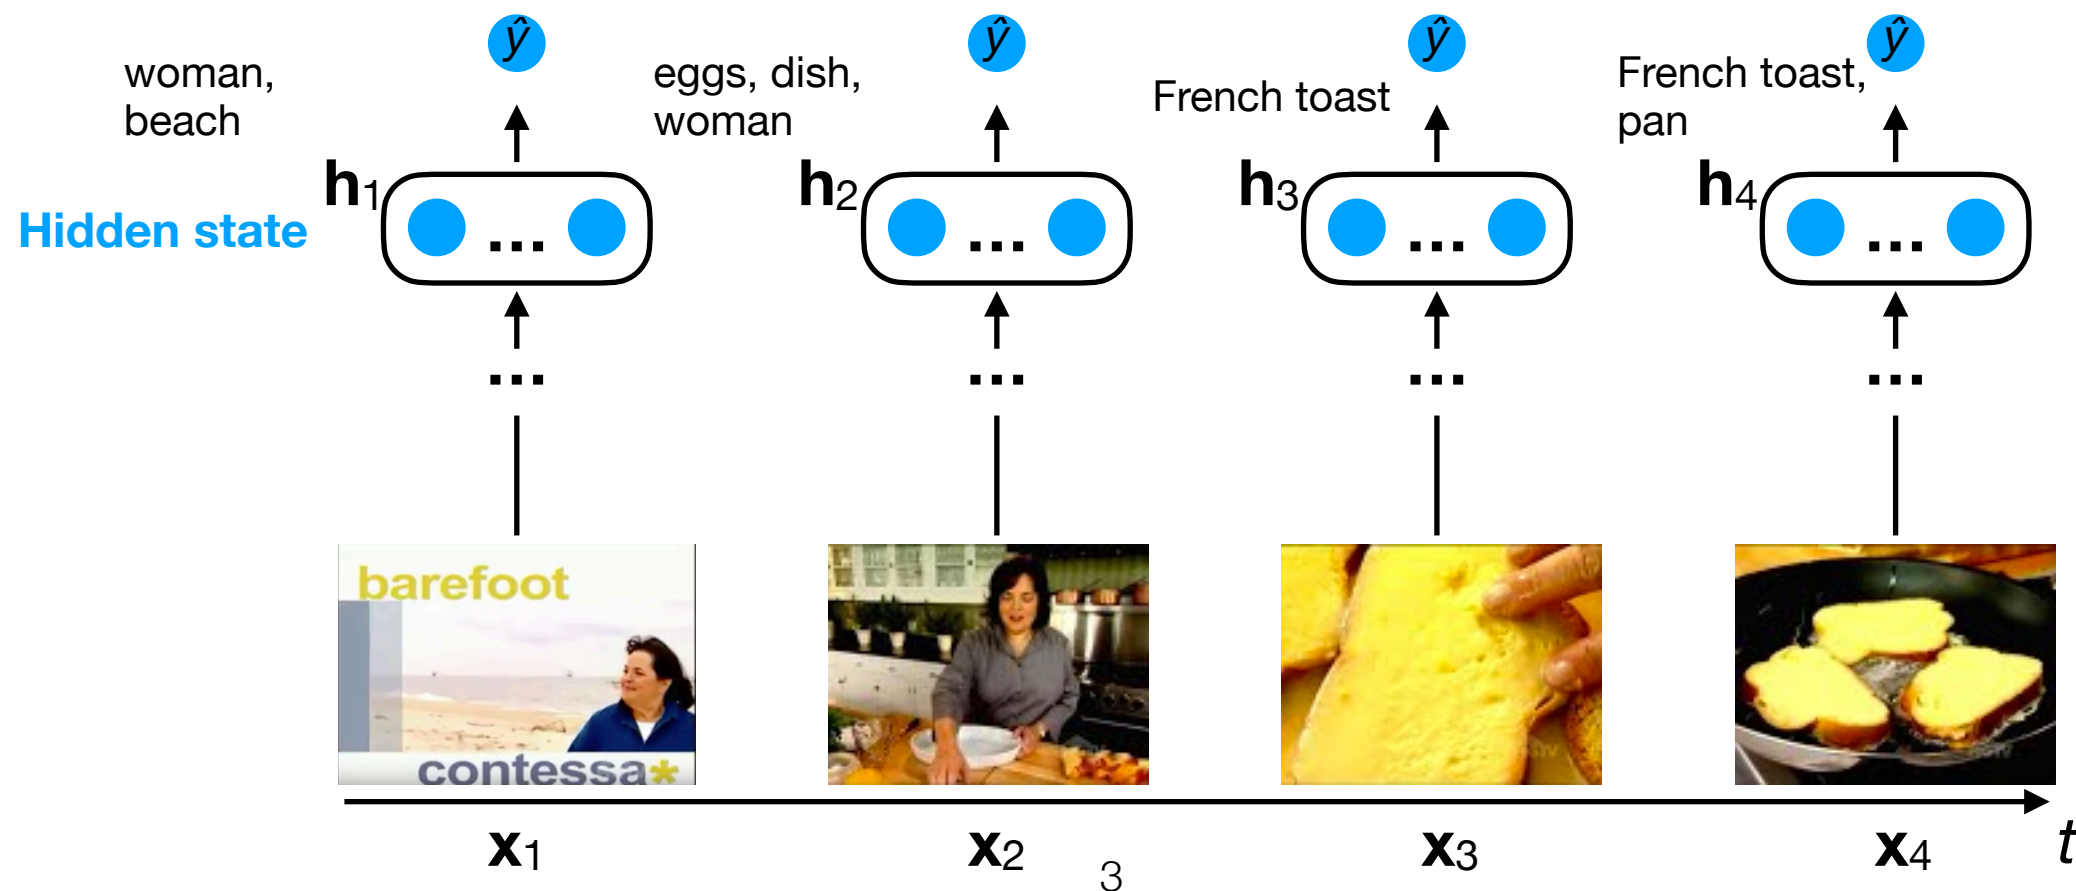
CS/DS 541: Class 12

Jacob Whitehill

Recurrent neural networks (RNNs)

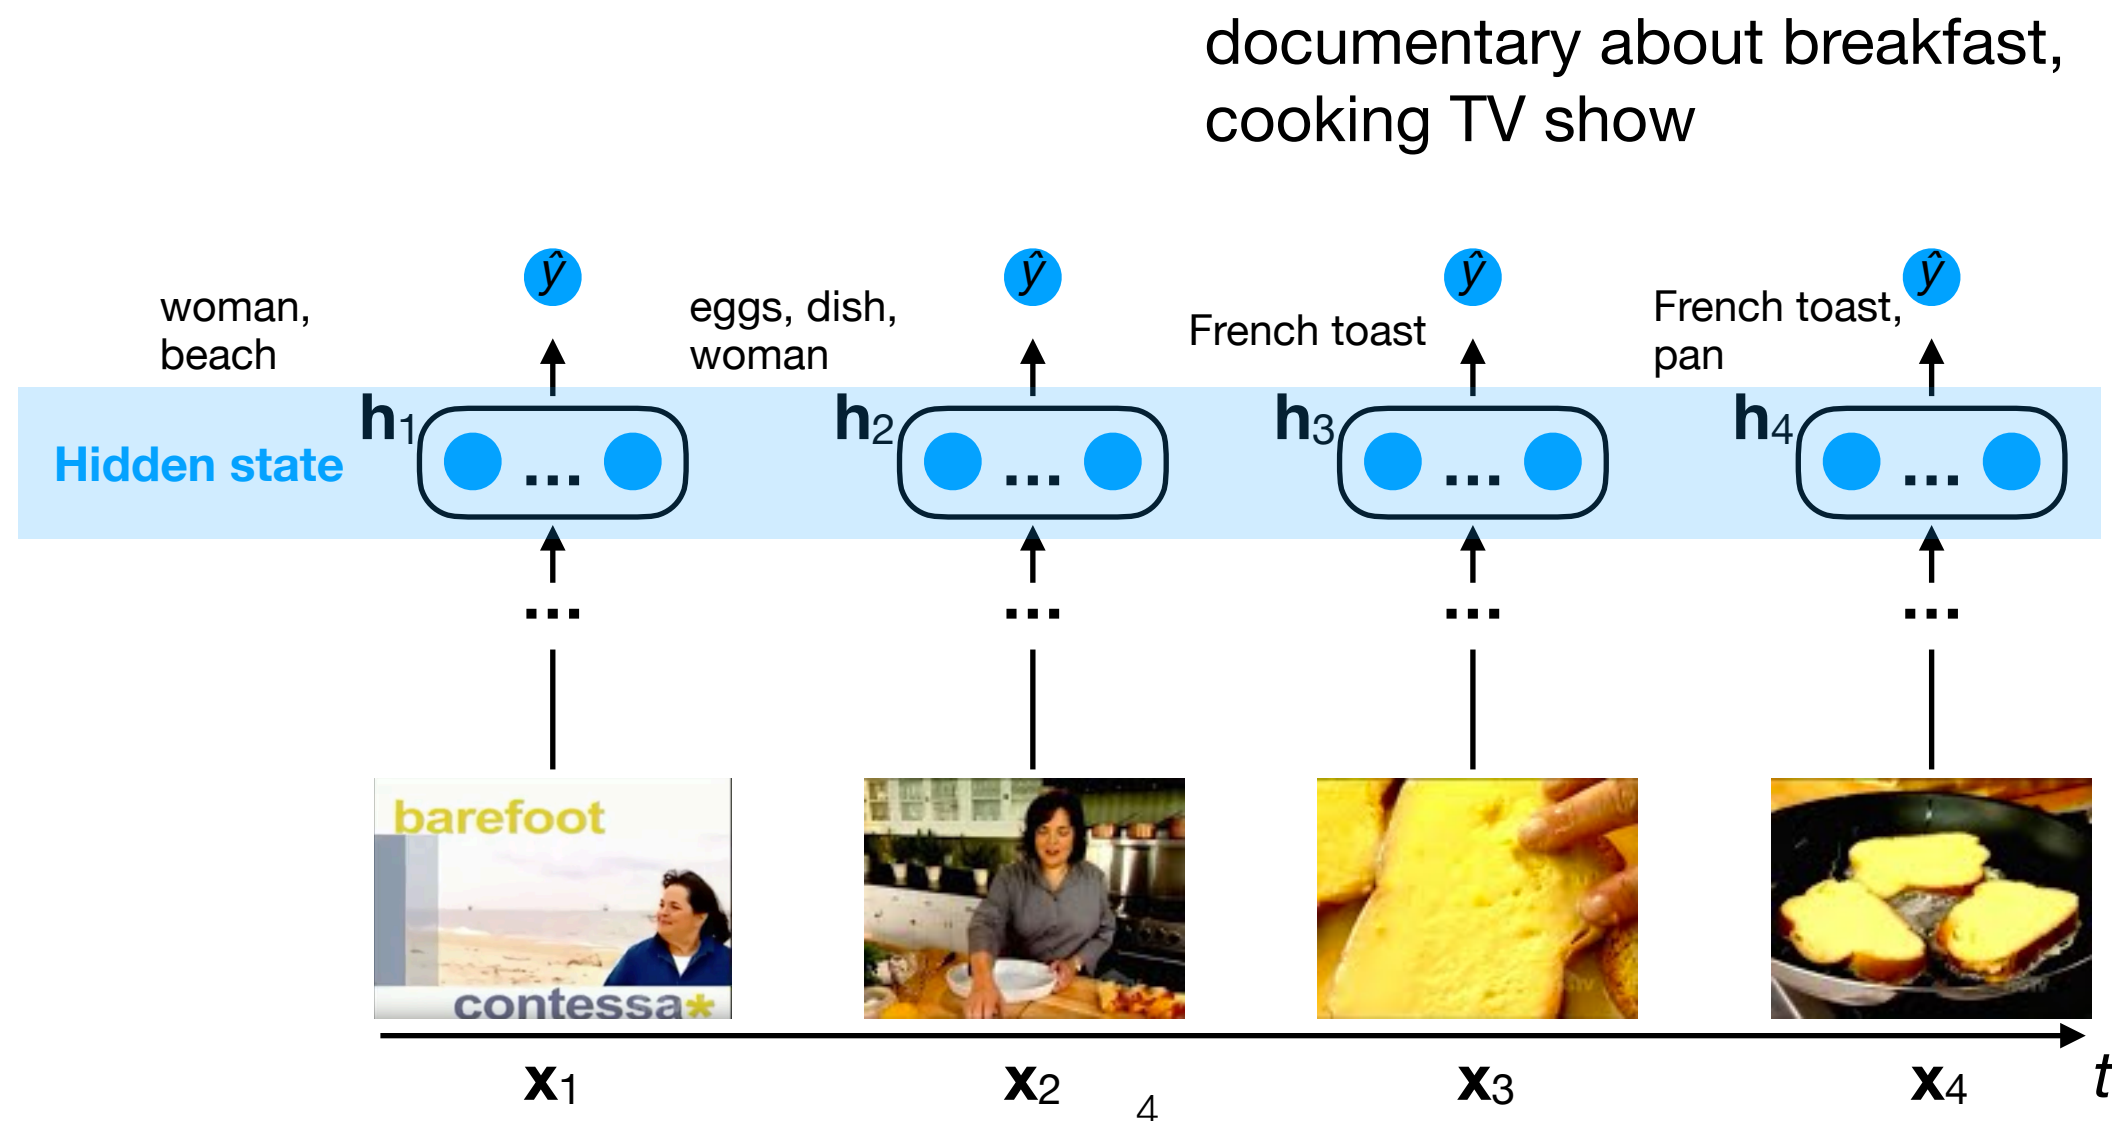
Role of the hidden state

- A CNN applied to each image in the sequence can tell us about the *individual objects* contained in *each frame*.
- Here, the hidden state \mathbf{h}_t of each input \mathbf{x}_t is computed **independently**, i.e.: $\mathbf{h}_t = f(\mathbf{x}_t)$



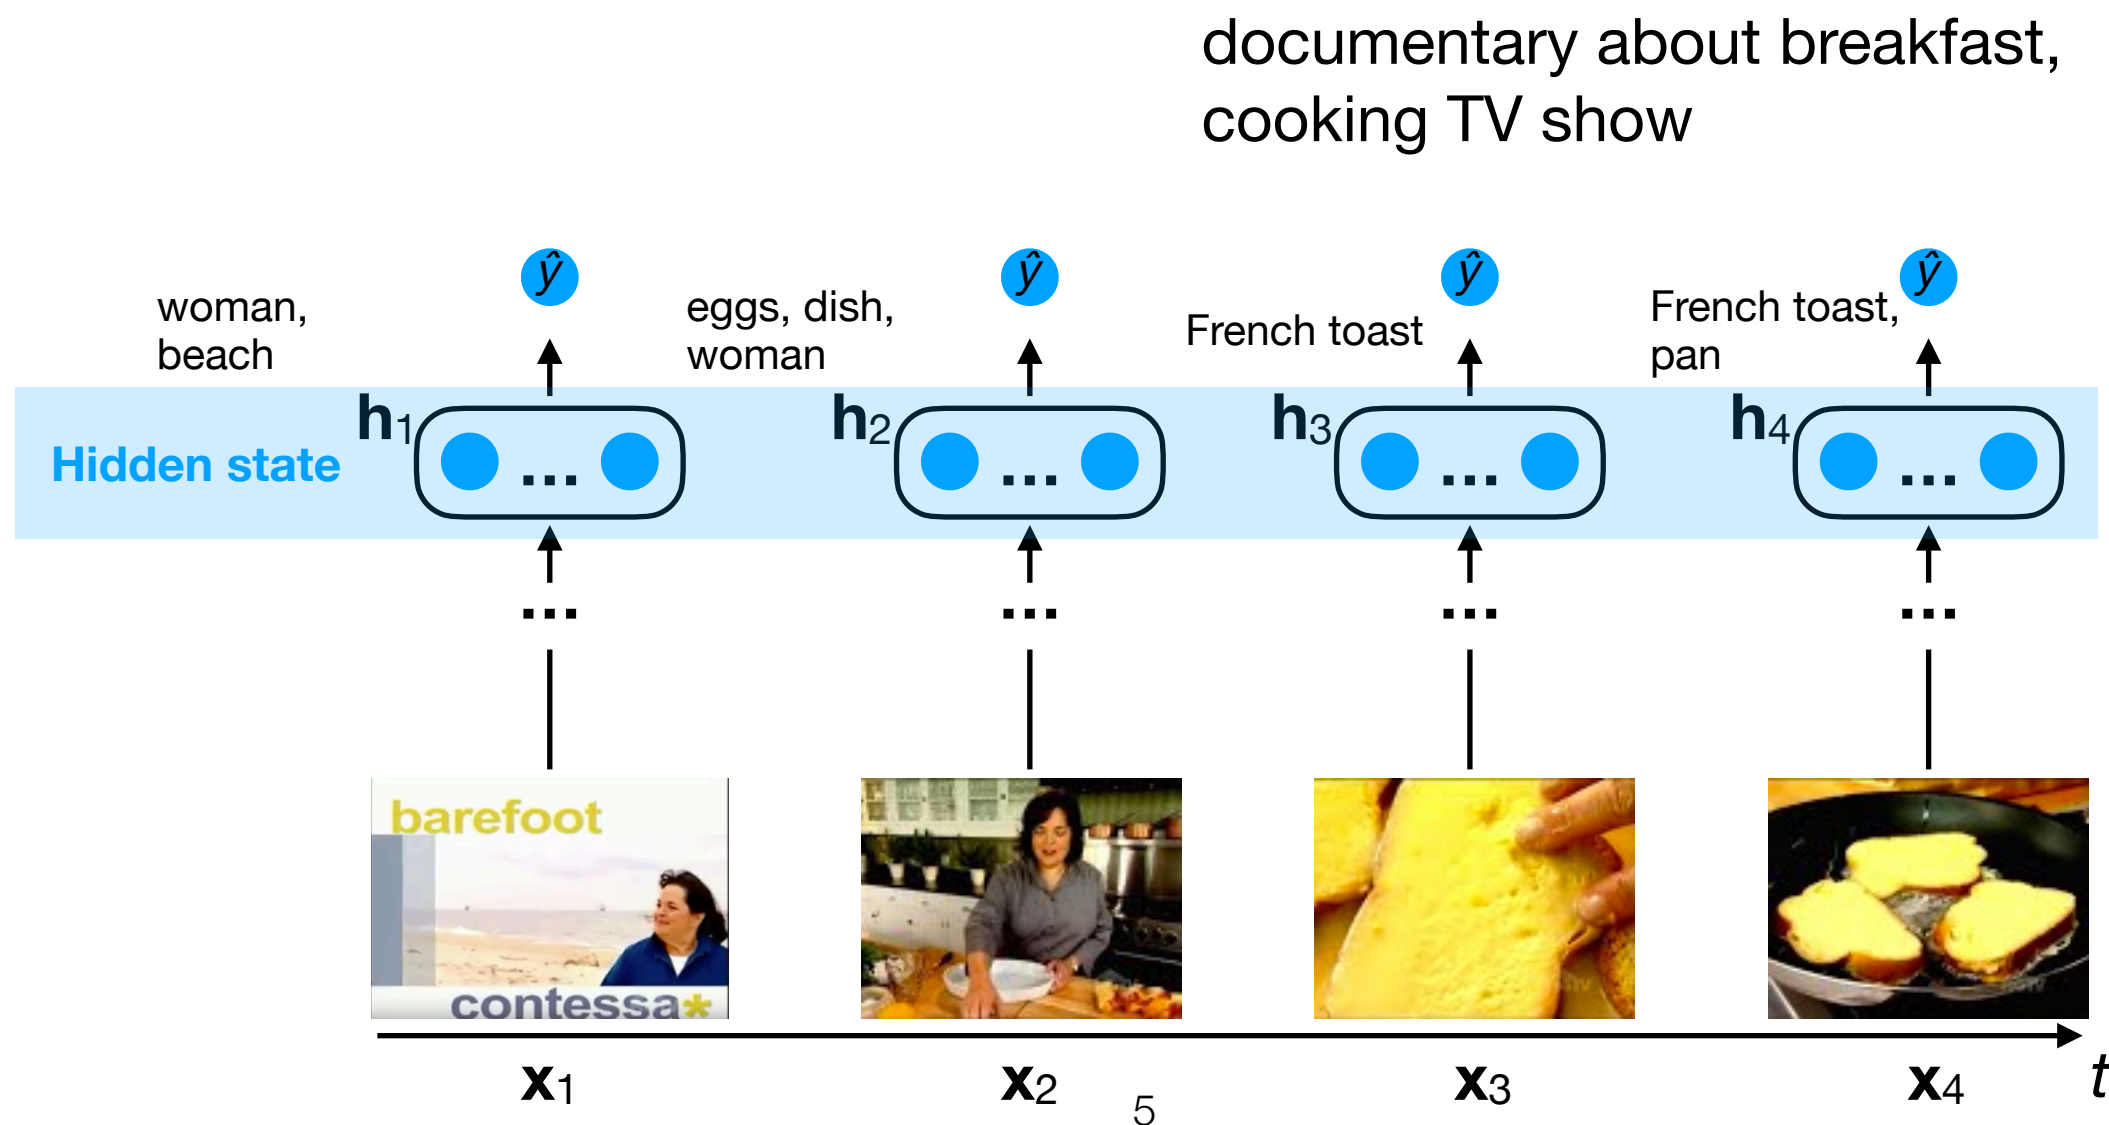
Role of the hidden state

- However, sometimes we need to **aggregate over time** to understand how the individual objects relate to each other:



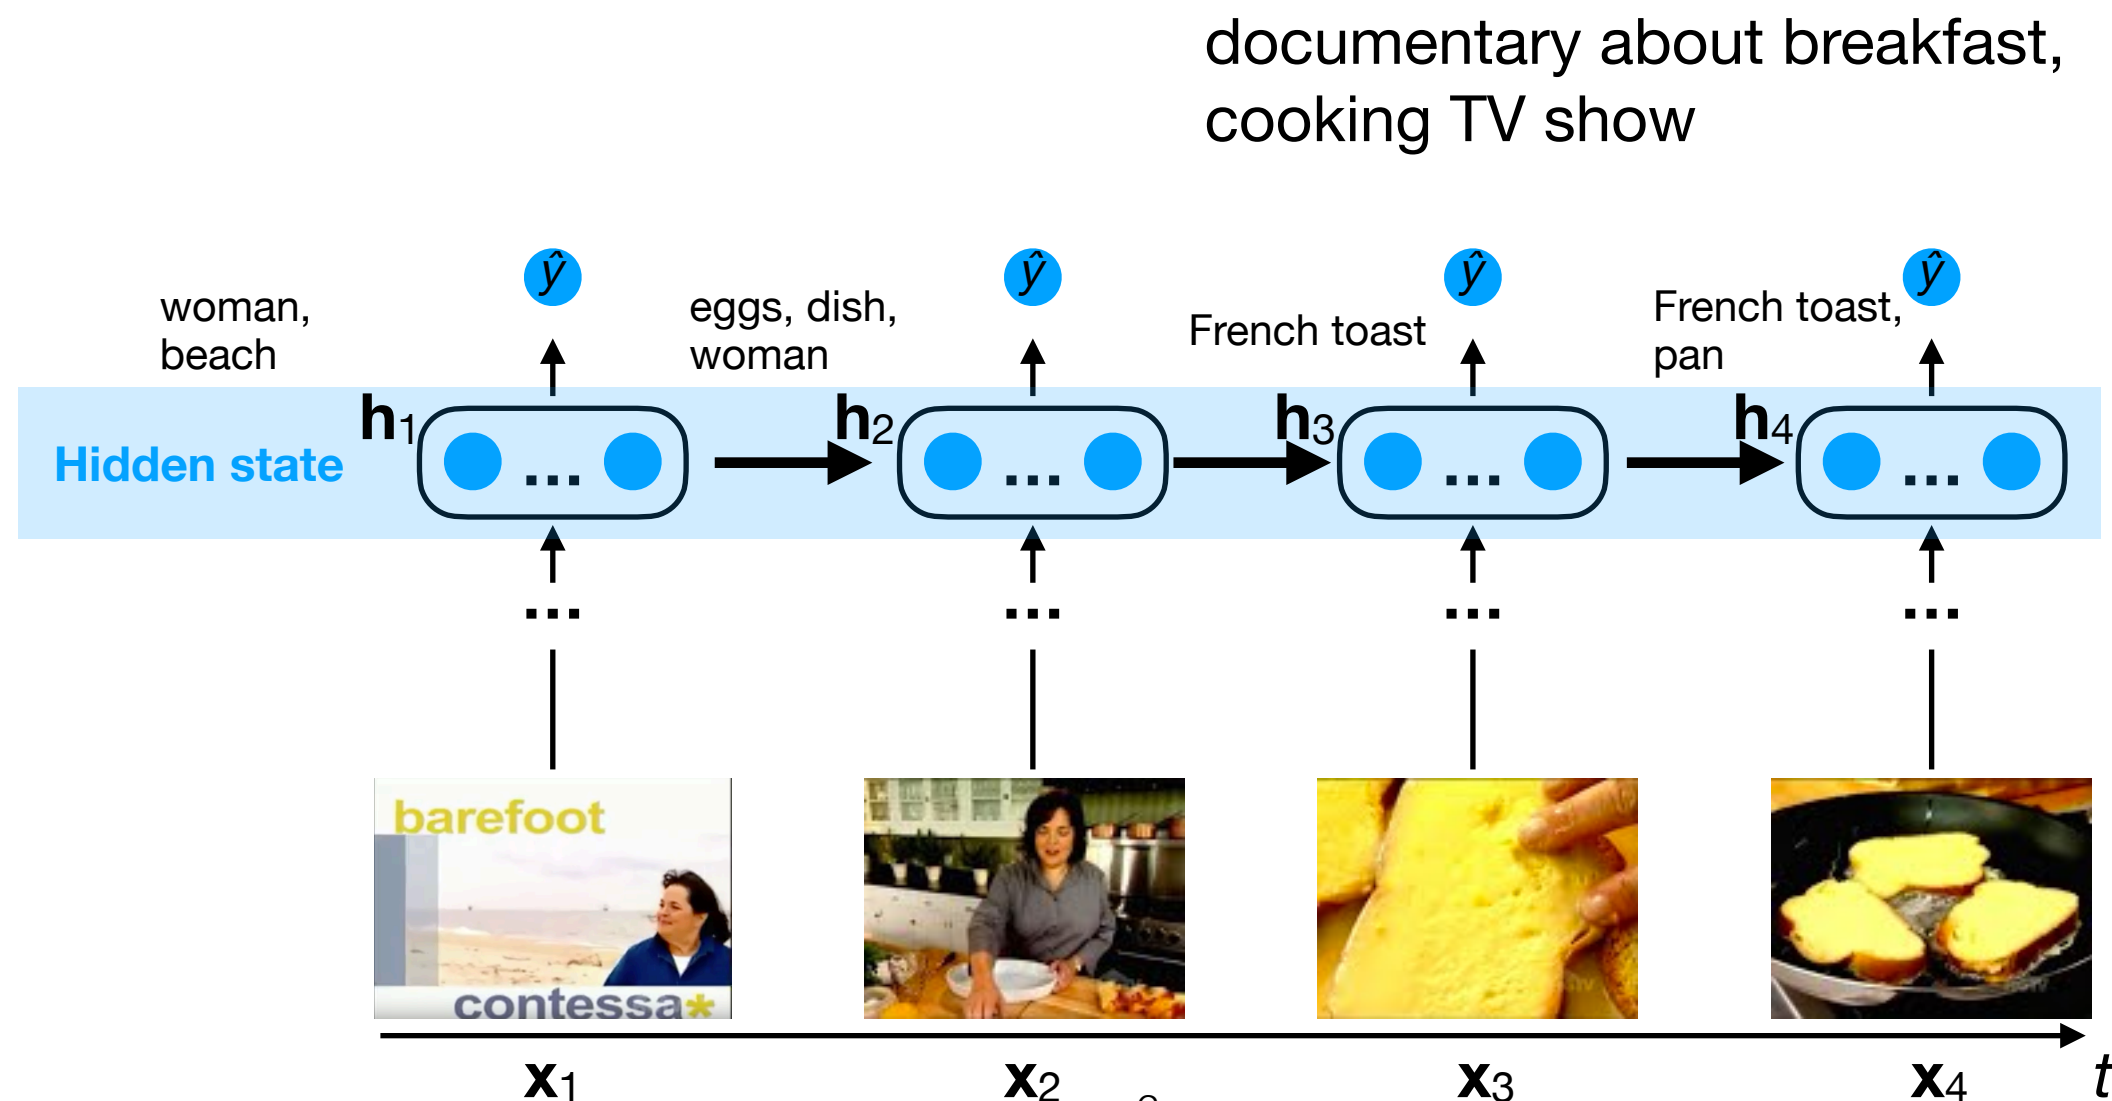
Role of the hidden state

- It's also possible that past hidden state helps us to infer the current hidden state (e.g., the French toast looks blurry at time $t=4$ but was very clear at time $t=3$).



Role of the hidden state

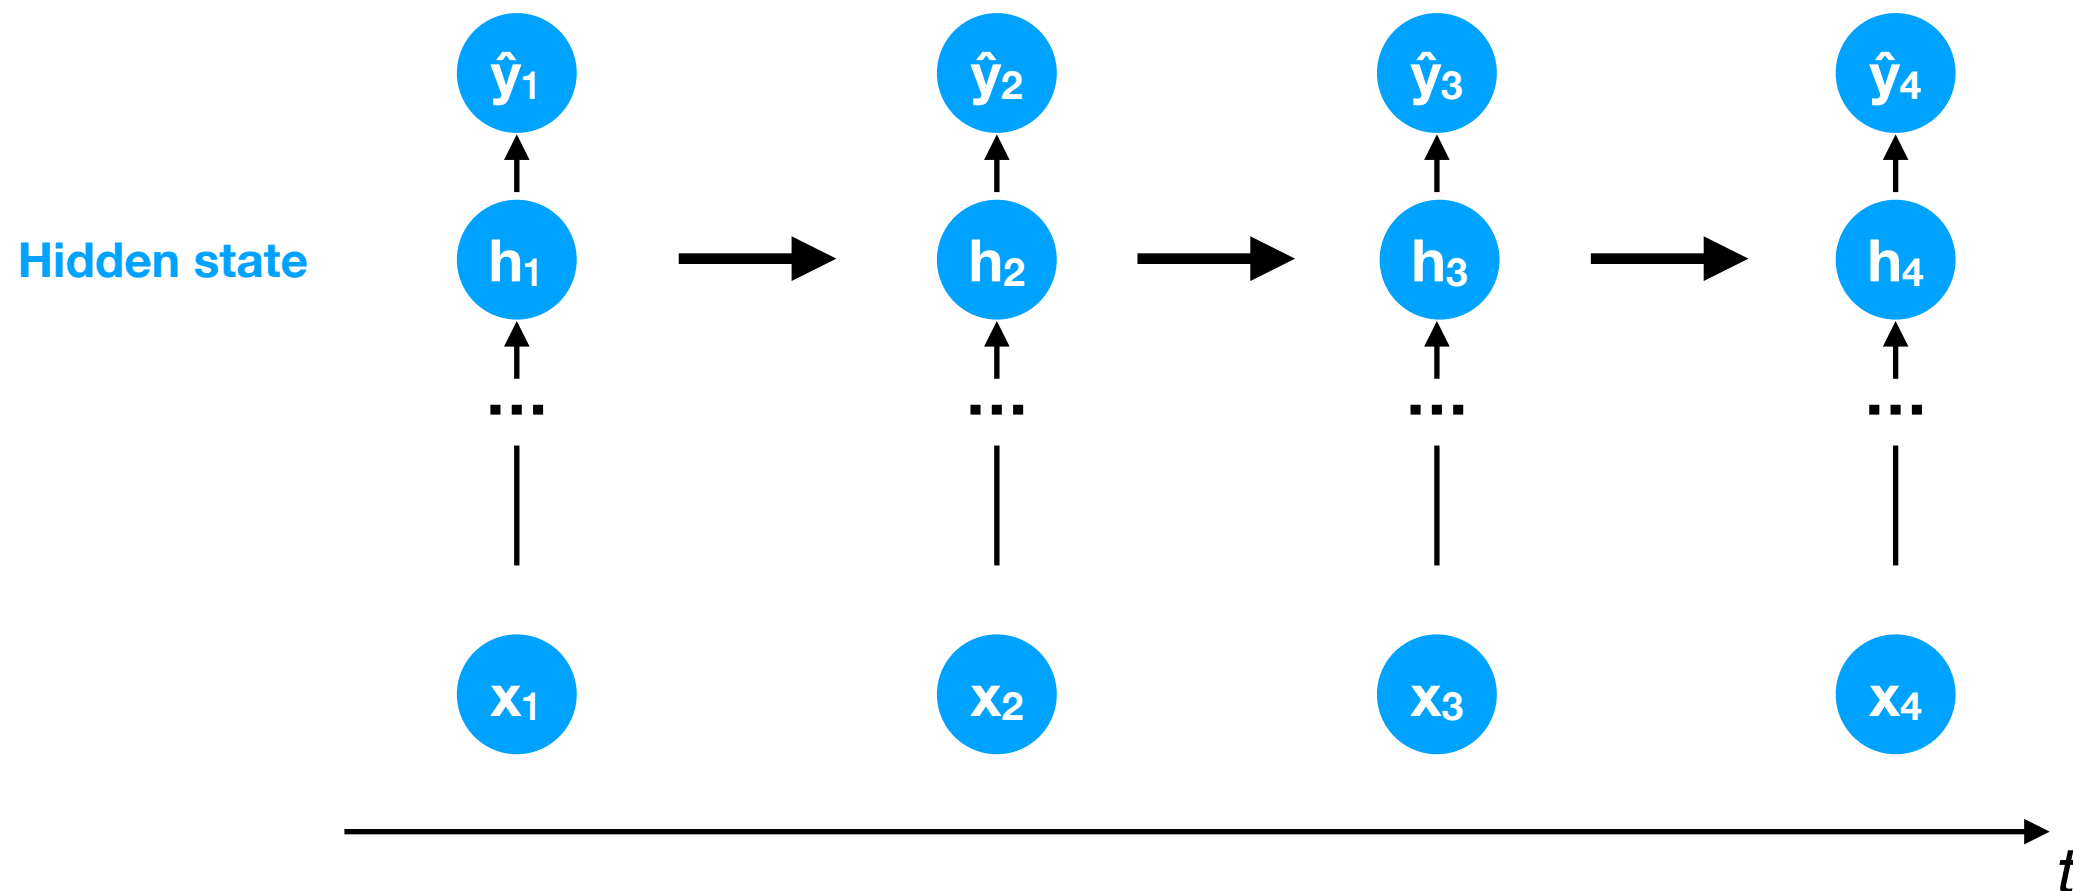
- To accomplish this, we can link the hidden state across the elements of the sequence, i.e.: $\mathbf{h}_t = f(\mathbf{x}_t, \mathbf{h}_{t-1})$



Recurrent neural network

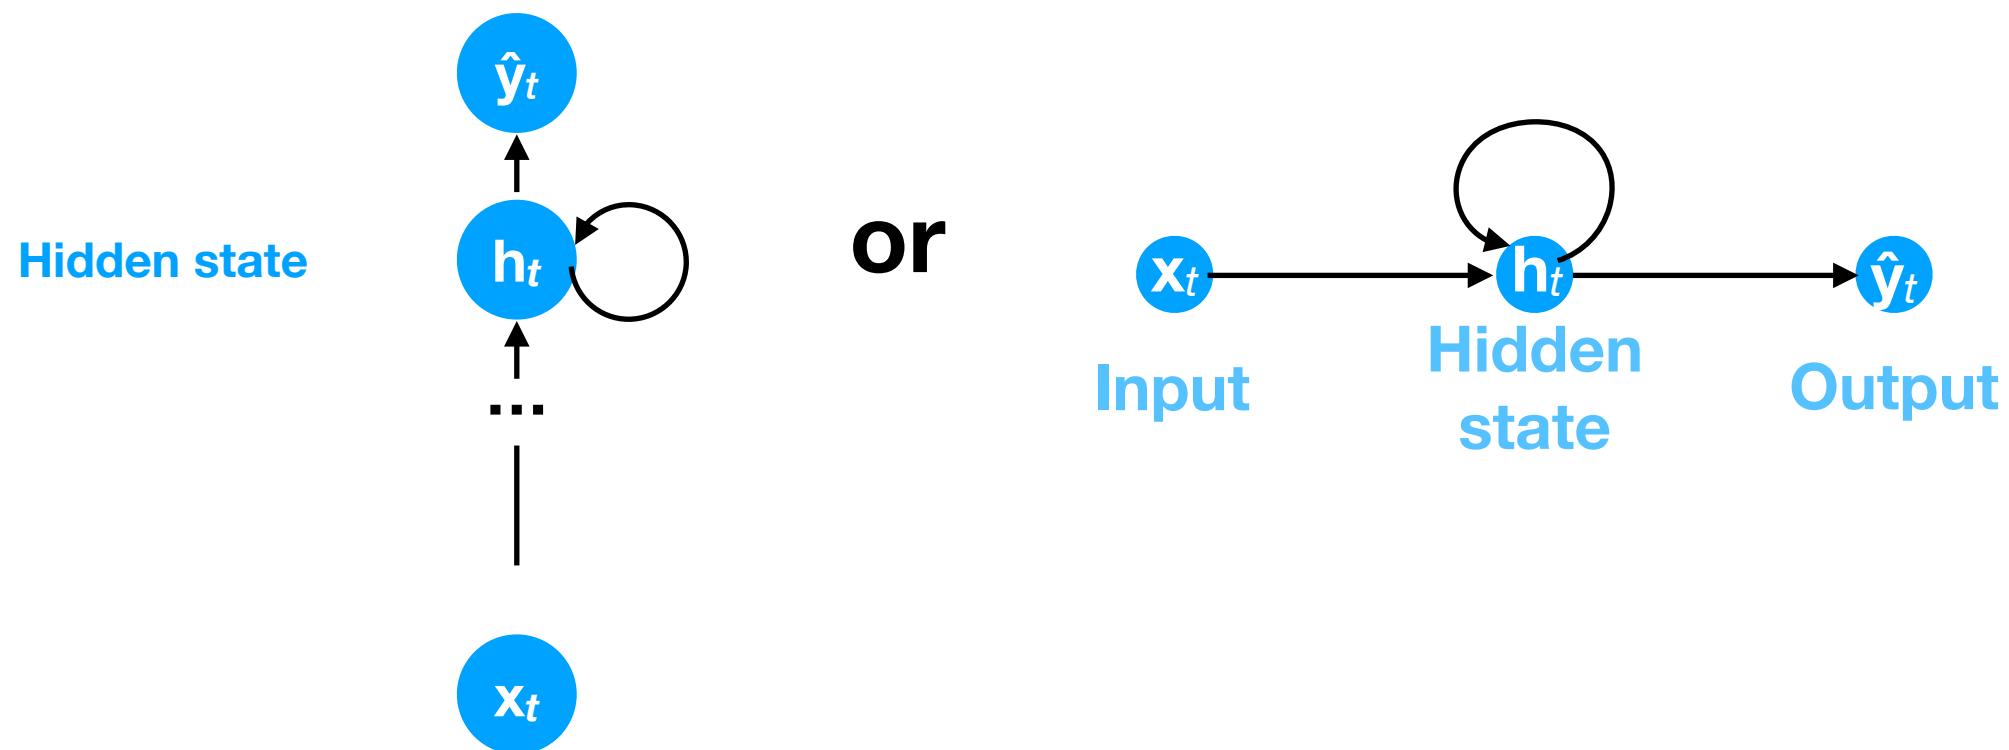
- This is the essence of a recurrent neural network (RNN) — the hidden states at time t depend on the hidden states of time $t-1$:

$$\mathbf{h}_t = f(\mathbf{x}_t, \mathbf{h}_{t-1})$$




Recurrent neural network

- We sometimes represent the recurrent hidden state as a single node with an arrow to itself:



Recurrent neural network

- We can construct a simple **recurrent neural network** (RNN) as follows:



$$\begin{aligned}\hat{y}_t &= g(\mathbf{x}_1, \dots, \mathbf{x}_t; \mathbf{U}, \mathbf{V}, \mathbf{w}) &= \mathbf{h}_t^\top \mathbf{w} \\ \mathbf{h}_t &= \sigma(\mathbf{U}\mathbf{h}_{t-1} + \mathbf{V}\mathbf{x}_t)\end{aligned}$$

Recurrent neural network

- As the activation σ , we typically use **tanh** instead of the logistic function:
 - Range of tanh: $(-1, +1)$
 - Range of logistic: $(0, 1)$
- Reason:
 - We want the hidden state to be able to maintain its value across many time steps ($t=1, 2, 3, \dots$).
 - $\tanh(x)$ is approximated by the line $y=x$ around $x=0$.

Recurrent neural network

- Consider a unidimensional hidden state. Let $h_1=0$.
- If $\sigma=\text{logistic}$, then:
 - $h_2=\sigma(h_1)=\sigma(0)=0.5$
 - $h_3=\sigma(h_2)=\sigma(0.5)=0.622$
 - $h_4=\sigma(h_3)=\sigma(0.622)=0.651$
 - ...

Recurrent neural network


- Consider a unidimensional hidden state. Let $h_1=0$.
- If $\sigma=\tanh$, then:
 - $h_2=\sigma(h_1)=\sigma(0)=0$
 - $h_3=\sigma(h_2)=\sigma(0)=0$
 - $h_4=\sigma(h_3)=\sigma(0)=0$
 - ...

Recurrent neural network

- The hidden state remains stable for tanh but not for logistic.
- Note that the *gradient* of tanh at 0 is exactly 1 because:
 $\tanh'(x) = 1 - \tanh^2(x)$
 - Learning can still occur!


Training a RNN

- We can use the same back-propagation procedure to train an RNN as we did for feed-forward neural networks (FFNN).
- For every input sequence $(\mathbf{x}_1, \dots, \mathbf{x}_T)$ and labels (y_1, \dots, y_T) , we **unroll** the computational graph over all T time-steps.




Training a RNN

- We can use the same back-propagation procedure to train an RNN as we did for feed-forward neural networks (FFNN).
- For every input sequence $(\mathbf{x}_1, \dots, \mathbf{x}_T)$ and labels (y_1, \dots, y_T) , we **unroll** the computational graph over all T time-steps.



Training a RNN


- The same weights (e.g., \mathbf{U}) are used in all time-steps, and we need to *sum* their contributions to the gradient $\nabla_{\mathbf{U}} J$ of the loss function J .



Training a RNN


- Moreover, we need to sum the loss values J_t at every timestep t :

$$J(\mathbf{U}) = \sum_{t=1}^T J_t(y_t, \hat{y}_t; \mathbf{U})$$




Training a RNN

- This process is called **back-propagation through time** (BPTT).



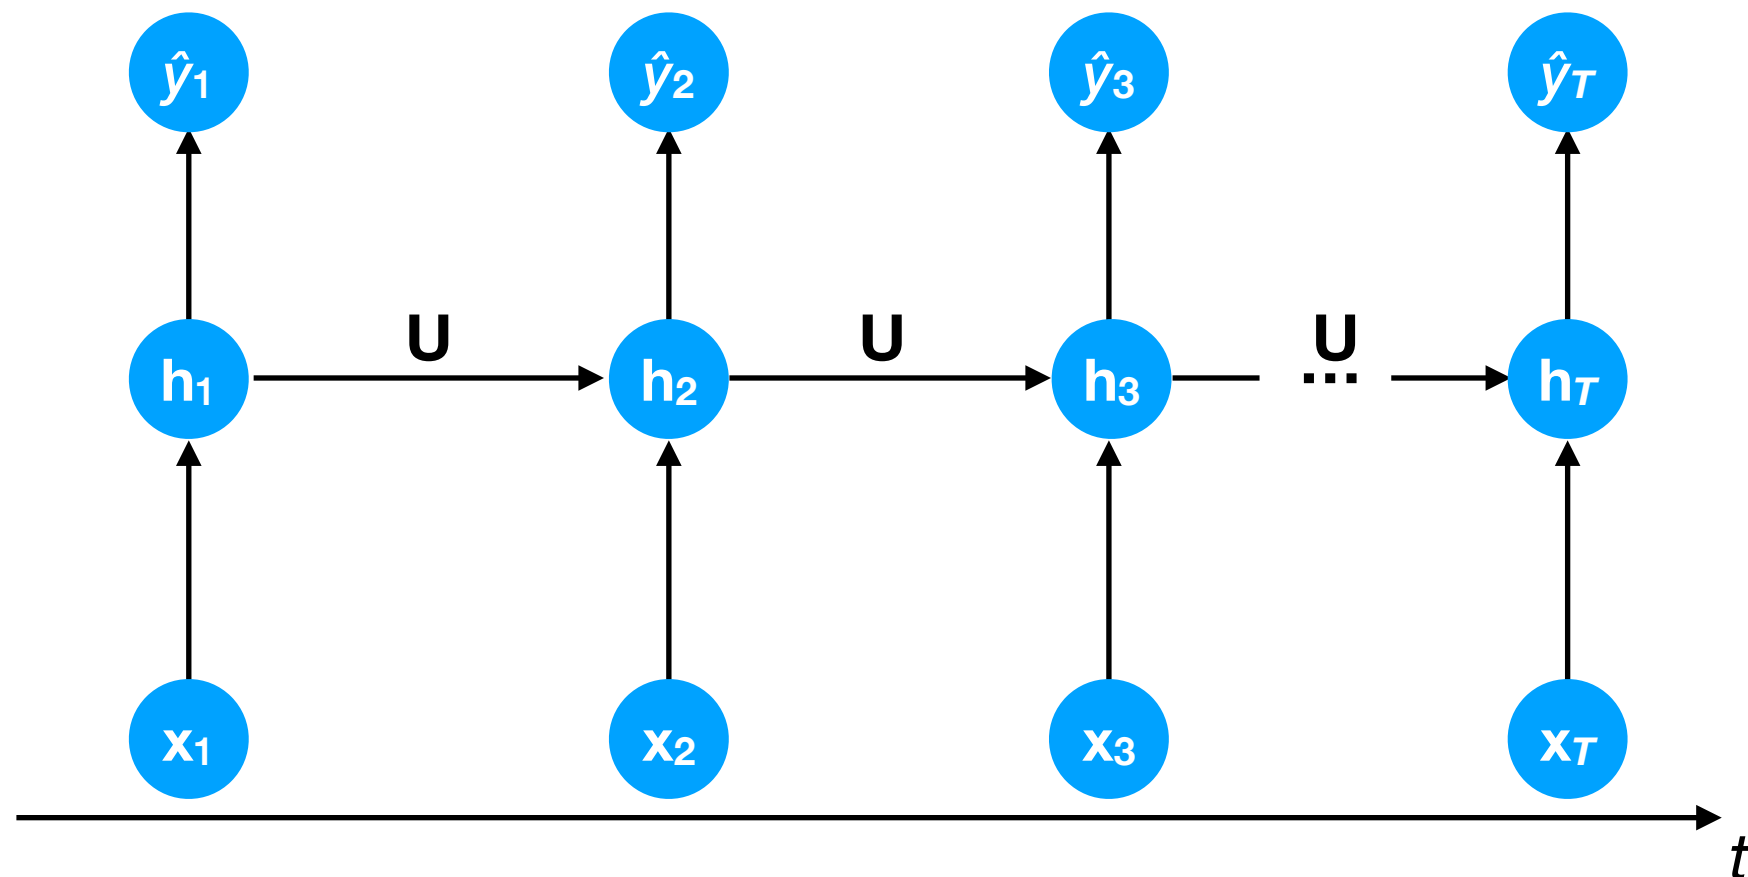
Depth in RNNs

- Note that, in the example RNN below, the “base” network is very simple and has just 1 hidden layer:




Depth in RNNs

- However, across timesteps, the maximum path-length from any input to any output can be large if T is large.
- In this sense, RNNs are deep networks.



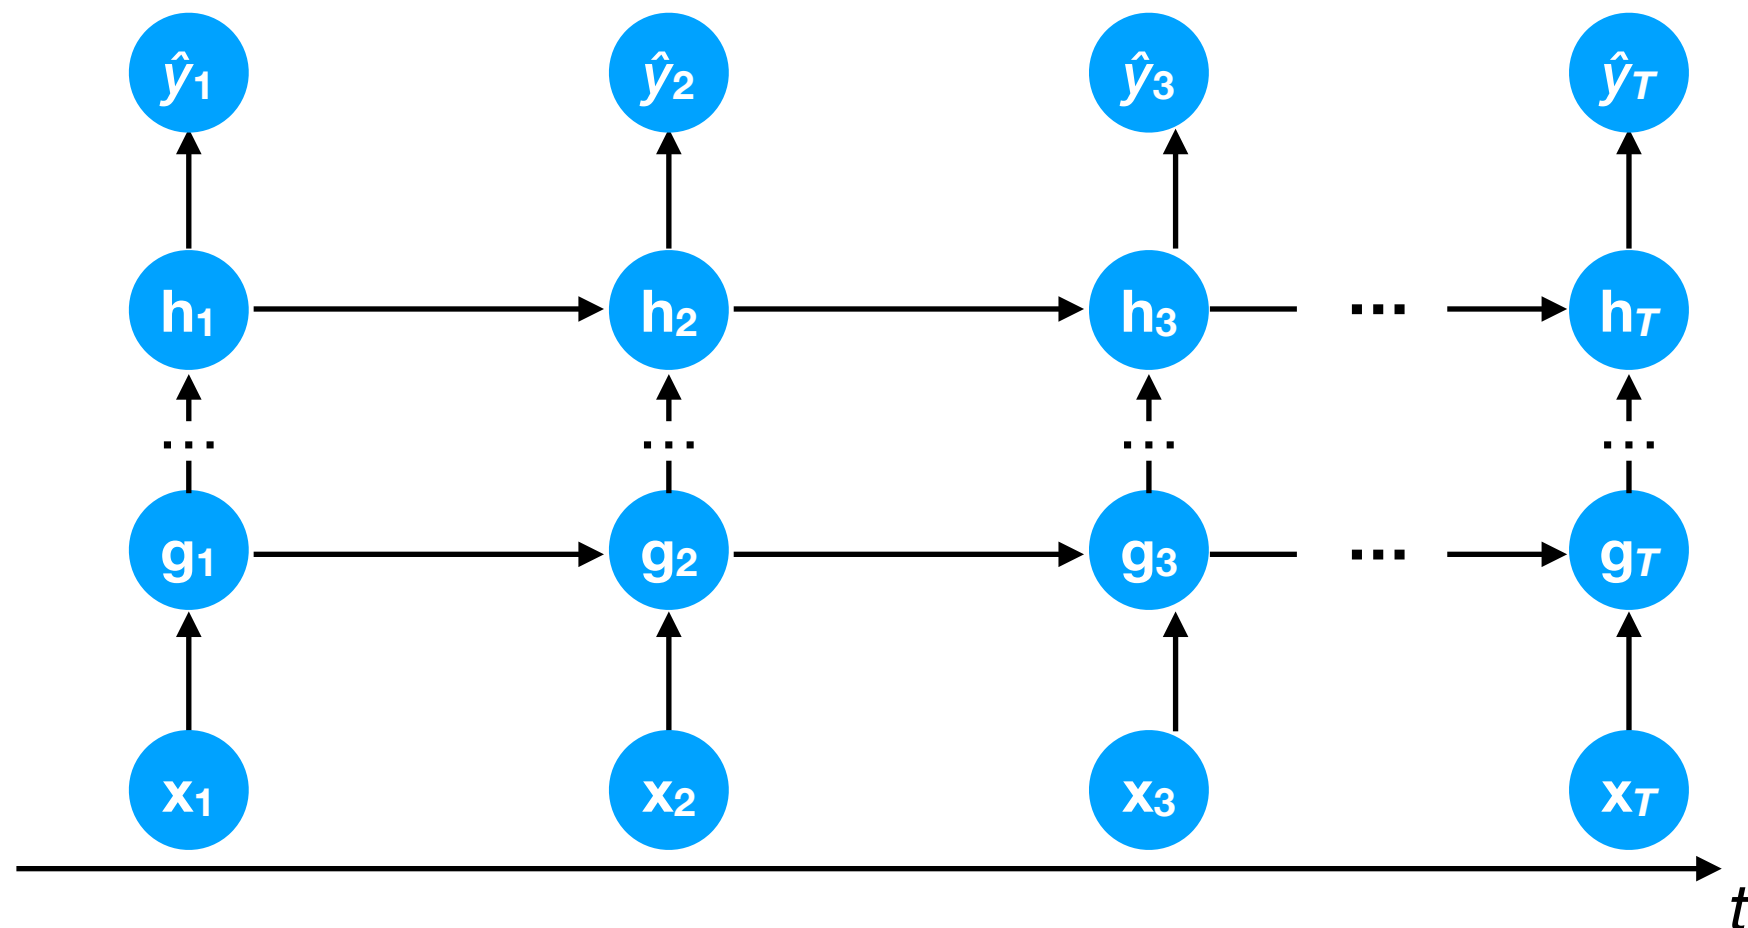
Depth in RNNs

- Note that we can also make the base RNN itself deep:
 - \mathbf{g}_t depends on \mathbf{x}_t and \mathbf{g}_{t-1} .
 - \mathbf{h}_t depends on \mathbf{g}_t and \mathbf{h}_{t-1} .



Depth in RNNs


- Note that we can also make the base RNN itself deeper:
 - \mathbf{g}_t depends on \mathbf{x}_t and \mathbf{g}_{t-1} .
 - \mathbf{h}_t depends on \mathbf{g}_t and \mathbf{h}_{t-1} .



Difficulty in training RNNs


Difficulty in training RNNs

- In their simplest form, RNNs are typically hard to train:
- The gradients can occasionally become very large (**exploding gradient**), which forces us to use a very small learning rate (which makes training slow).




Difficulty in training RNNs

- In their simplest form, RNNs are typically hard to train:
- The gradients can also become very small (**vanishing gradient**), which also makes learning very slow.



Difficulty in training RNNs


- A related problem is that, if T is large, then information *early* in the input sequence (e.g., \mathbf{x}_1) can “get lost” when trying to predict values *late* in the sequence (e.g., \hat{y}_T).



Difficulty in training RNNs

- Why do these problems occur?
- Let's consider a simplistic RNN *without* a non-linear activation function, i.e.:

$$\mathbf{h}_t = \mathbf{U}\mathbf{h}_{t-1} + \mathbf{V}\mathbf{x}_t$$




Difficulty in training RNNs

- When we compute the gradient of $J(y_T, \hat{y}_T)$ w.r.t. \mathbf{V} , we must sum over all occurrences of \mathbf{V} —including the *first* one:

$$\frac{\partial J_T(\mathbf{V})}{\partial \mathbf{V}} = \frac{\partial J_T(\mathbf{V})}{\partial \mathbf{h}_T} \frac{\partial \mathbf{h}_T}{\partial \mathbf{h}_{T-1}} \cdots \frac{\partial \mathbf{h}_2}{\partial \mathbf{h}_1} \frac{\partial \mathbf{h}_1}{\partial \mathbf{V}} + \text{other terms}$$


There are “other terms” because \mathbf{V} occurs at not just $t=1$.



Difficulty in training RNNs

- When we compute the gradient of $J(y_T, \hat{y}_T)$ w.r.t. \mathbf{V} , we must sum over all occurrences of \mathbf{V} —including the *first* one:

$$\frac{\partial J_T(\mathbf{V})}{\partial \mathbf{V}} = \frac{\partial J_T(\mathbf{V})}{\partial \mathbf{h}_T} \mathbf{U} \mathbf{U} \dots \mathbf{U} \frac{\partial \mathbf{h}_1}{\partial \mathbf{V}}$$



Difficulty in training RNNs

- Note that **U** is square (it maps from one hidden state to another).
- If **U** is diagonalizable, then we can write it as:

$$\mathbf{U} = \mathbf{Q}\mathbf{D}\mathbf{Q}^{-1}$$

where the columns of **Q** are **U**'s eigenvectors and **D** contains **U**'s the associated eigenvalues.

Difficulty in training RNNs

- Note that **U** is square (it maps from one hidden state to another).
- If **U** is diagonalizable, then we can write it as:

$$\mathbf{U} = \mathbf{Q}\mathbf{D}\mathbf{Q}^{-1}$$

where the columns of **Q** are **U**'s eigenvectors and **D** contains **U**'s the associated eigenvalues.

- Hence:

$$\overbrace{\mathbf{U}\mathbf{U} \dots \mathbf{U}}^{k \text{ terms}} =$$

Difficulty in training RNNs

- Note that **U** is square (it maps from one hidden state to another).
- If **U** is diagonalizable, then we can write it as:

$$\mathbf{U} = \mathbf{Q}\mathbf{D}\mathbf{Q}^{-1}$$

where the columns of **Q** are **U**'s eigenvectors and **D** contains **U**'s the associated eigenvalues.

- Hence:

$$\begin{aligned}\overbrace{\mathbf{U}\mathbf{U} \dots \mathbf{U}}^{k \text{ terms}} &= \mathbf{Q}\mathbf{D}\mathbf{Q}^{-1}\mathbf{Q}\mathbf{D}\mathbf{Q}^{-1} \dots \mathbf{Q}\mathbf{D}\mathbf{Q}^{-1} \\ &= \mathbf{Q}\mathbf{D}^k\mathbf{Q}^{-1}\end{aligned}$$

Difficulty in training RNNs

- Plugging back into the gradient:

$$\begin{aligned}\frac{\partial J_T(\mathbf{V})}{\partial \mathbf{V}} &= \frac{\partial J_T(\mathbf{V})}{\partial \mathbf{h}_T} \mathbf{U} \mathbf{U} \dots \mathbf{U} \frac{\partial \mathbf{h}_1}{\partial \mathbf{V}} \\ &= \frac{\partial J_T(\mathbf{V})}{\partial \mathbf{h}_T} \mathbf{Q} \mathbf{D}^{T-1} \mathbf{Q}^{-1} \frac{\partial \mathbf{h}_1}{\partial \mathbf{V}} \\ &= \frac{\partial J_T(\mathbf{V})}{\partial \mathbf{h}_T} \mathbf{Q} \begin{bmatrix} \lambda_1^{T-1} & \dots & 0 \\ 0 & \ddots & 0 \\ 0 & \dots & \lambda_m^{T-1} \end{bmatrix} \mathbf{Q}^{-1} \frac{\partial \mathbf{h}_1}{\partial \mathbf{V}}\end{aligned}$$

Difficulty in training RNNs

- Plugging back into the gradient:

$$\begin{aligned}
 \frac{\partial J_T(\mathbf{V})}{\partial \mathbf{V}} &= \frac{\partial J_T(\mathbf{V})}{\partial \mathbf{h}_T} \mathbf{U} \mathbf{U} \dots \mathbf{U} \frac{\partial \mathbf{h}_1}{\partial \mathbf{V}} \\
 &= \frac{\partial J_T(\mathbf{V})}{\partial \mathbf{h}_T} \mathbf{Q} \mathbf{D}^{T-1} \mathbf{Q}^{-1} \frac{\partial \mathbf{h}_1}{\partial \mathbf{V}} \\
 &= \frac{\partial J_T(\mathbf{V})}{\partial \mathbf{h}_T} \mathbf{Q} \begin{bmatrix} \lambda_1^{T-1} & \dots & 0 \\ 0 & \ddots & 0 \\ 0 & \dots & \lambda_m^{T-1} \end{bmatrix} \mathbf{Q}^{-1} \frac{\partial \mathbf{h}_1}{\partial \mathbf{V}}
 \end{aligned}$$

- If $|\lambda_i| < 1$, then $\lambda_i^{T-1} \rightarrow 0$ as $T \rightarrow \infty$. **Vanishing gradient.**
- If $|\lambda_i| > 1$, then $|\lambda_i^{T-1}| \rightarrow \infty$ as $T \rightarrow \infty$. **Exploding gradient.**

Difficulty in training RNNs

- Moreover, during forward-propagation:

$$\hat{y}_T = \mathbf{w} \mathbf{U} \mathbf{U} \dots \mathbf{U} \mathbf{x}_1 + \text{other terms}$$

$$= \mathbf{w} \mathbf{Q} \mathbf{D}^{T-1} \mathbf{Q}^{-1} \mathbf{x}_1$$

$$= \mathbf{w} \begin{bmatrix} \mathbf{u}_1 & \dots & \mathbf{u}_m \end{bmatrix} \begin{bmatrix} \lambda_1^{T-1} & \dots & 0 \\ 0 & \ddots & 0 \\ 0 & \dots & \lambda_m^{T-1} \end{bmatrix} \begin{bmatrix} \mathbf{u}_1^\top \\ \vdots \\ \mathbf{u}_m^\top \end{bmatrix} \mathbf{x}_1$$

where \mathbf{u}_i is the i^{th} eigenvector of \mathbf{U} .

- \hat{y}_T loses information from \mathbf{x}_1 along direction \mathbf{u}_i unless $|\lambda_i| > 1$.

Difficulty in training deep FFNNs

- Note that the same vanishing/exploding gradient problems can also occur in very deep FFNNs, particularly if the rows of the weight matrices are highly correlated.

Difficulty in training deep FFNNs

- One coping strategy is to **clip** the gradients:
- Suppose we obtain a gradient $\nabla_p J$ for some parameter p , and $|\nabla_p J|$ is very large (above some threshold τ).
- Then we can “clip” its value to be:

$$\nabla_p J \left(\frac{\tau}{|\nabla_p J|} \right)$$

Difficulty in training deep FFNNs

- Another strategy for preventing vanishing and exploding gradients is to use **skip connections** (more later).
 - These are used in LSTM and GRU RNNs, as well as ResNet FFNNs.
- Yet another strategy is to restrict **U** to the manifold of **unitary matrices** (i.e., all eigenvalues have magnitude 1; see Helfrich & Ye 2019).

Training RNNs

- Like CNNs, RNNs exhibit an inherently large degree of weight sharing:
 - The **U**, **V** matrices are the same at *all* timesteps.
- See rnn.pdf on Canvas.

Long short-term memory (LSTM) neural networks

LSTMs

- <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

LSTMs

- Three gates — forget (f), input (i), and output (o).

$$\mathbf{f}_t = \sigma(\mathbf{W}_f[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_f)$$

$$\mathbf{i}_t = \sigma(\mathbf{W}_i[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_i)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_o)$$

LSTMs

- Three gates — forget (f), input (i), and output (o).
- Two state vectors: \mathbf{h}_t , \mathbf{c}_t .

$$\mathbf{f}_t = \sigma(\mathbf{W}_f[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_f)$$

$$\mathbf{i}_t = \sigma(\mathbf{W}_i[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_i)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_o)$$

$$\tilde{\mathbf{c}}_t = \tanh(\mathbf{W}_c[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_c)$$

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$$

LSTMs

- In total, we have 4 weight matrices and 4 bias vectors.

$$\mathbf{f}_t = \sigma(\mathbf{W}_f[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_f)$$

$$\mathbf{i}_t = \sigma(\mathbf{W}_i[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_i)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_o)$$

$$\tilde{\mathbf{c}}_t = \tanh(\mathbf{W}_c[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_c)$$

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t$$


$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$$

LSTM

- The memory cell \mathbf{c} offers a pathway through the network to preserve information across long time-spans:

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{f}_t \odot \mathbf{c}_{t-1}$$


- It tends not to decay due to exponentiated eigenvalues.



LSTM

- If $\mathbf{f}_t=1$, then \mathbf{c}_t directly contains information from 1, ..., t :

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1}$$




LSTM

- If $\mathbf{f}_t=1$, then \mathbf{c}_t directly contains information from 1, ..., t :

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{c}_{t-2}$$



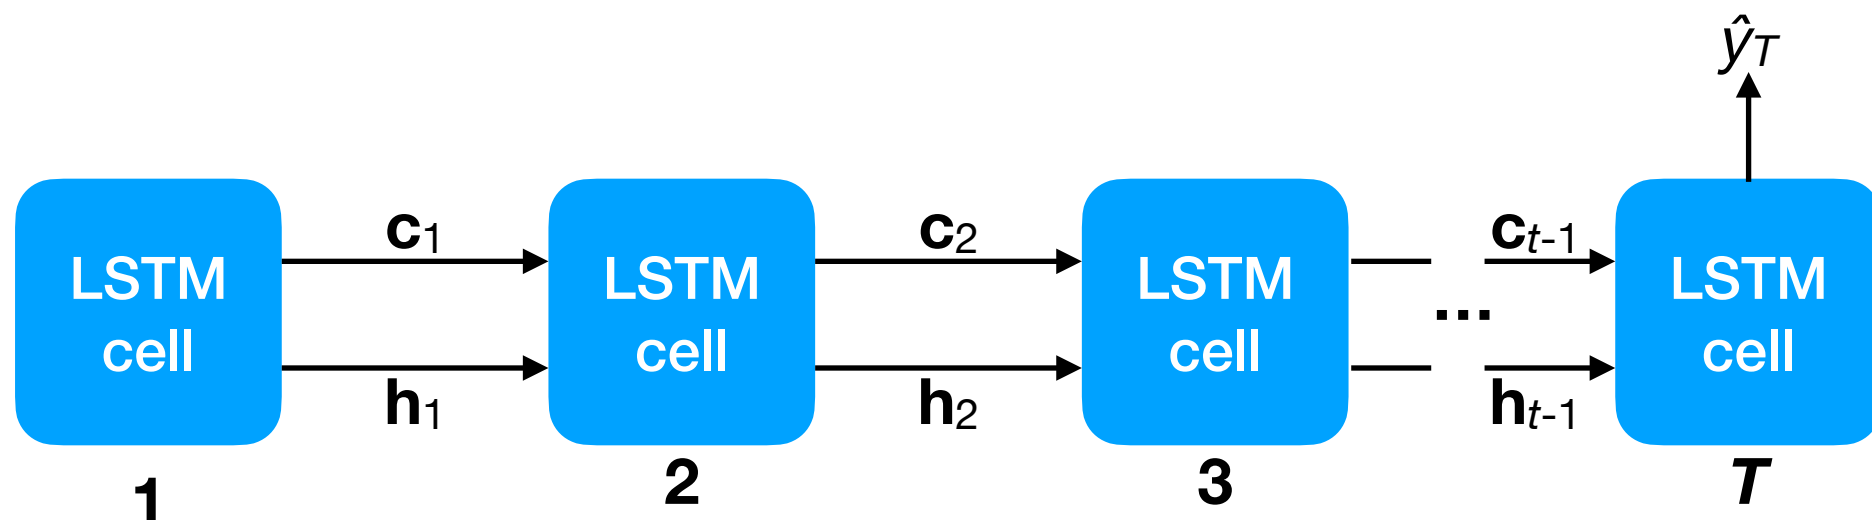
LSTM

- If $\mathbf{f}_t=1$, then \mathbf{c}_t directly contains information from 1, ..., t :

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{c}_{t-2}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{i}_{t-2} \odot \tilde{\mathbf{c}}_{t-2} + \mathbf{c}_{t-3}$$



LSTM

- If $\mathbf{f}_t=1$, then \mathbf{c}_t directly contains information from 1, ..., t :


$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{c}_{t-2}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{i}_{t-2} \odot \tilde{\mathbf{c}}_{t-2} + \mathbf{c}_{t-3}$$

...

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \dots + \mathbf{i}_2 \odot \tilde{\mathbf{c}}_2 + \mathbf{c}_1$$



LSTM

- If $\mathbf{f}_t=1$, then \mathbf{c}_t directly contains information from 1, ..., t :

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{c}_{t-2}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{i}_{t-2} \odot \tilde{\mathbf{c}}_{t-2} + \mathbf{c}_{t-3}$$

...

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \dots + \mathbf{i}_2 \odot \tilde{\mathbf{c}}_2 + \mathbf{c}_1$$

